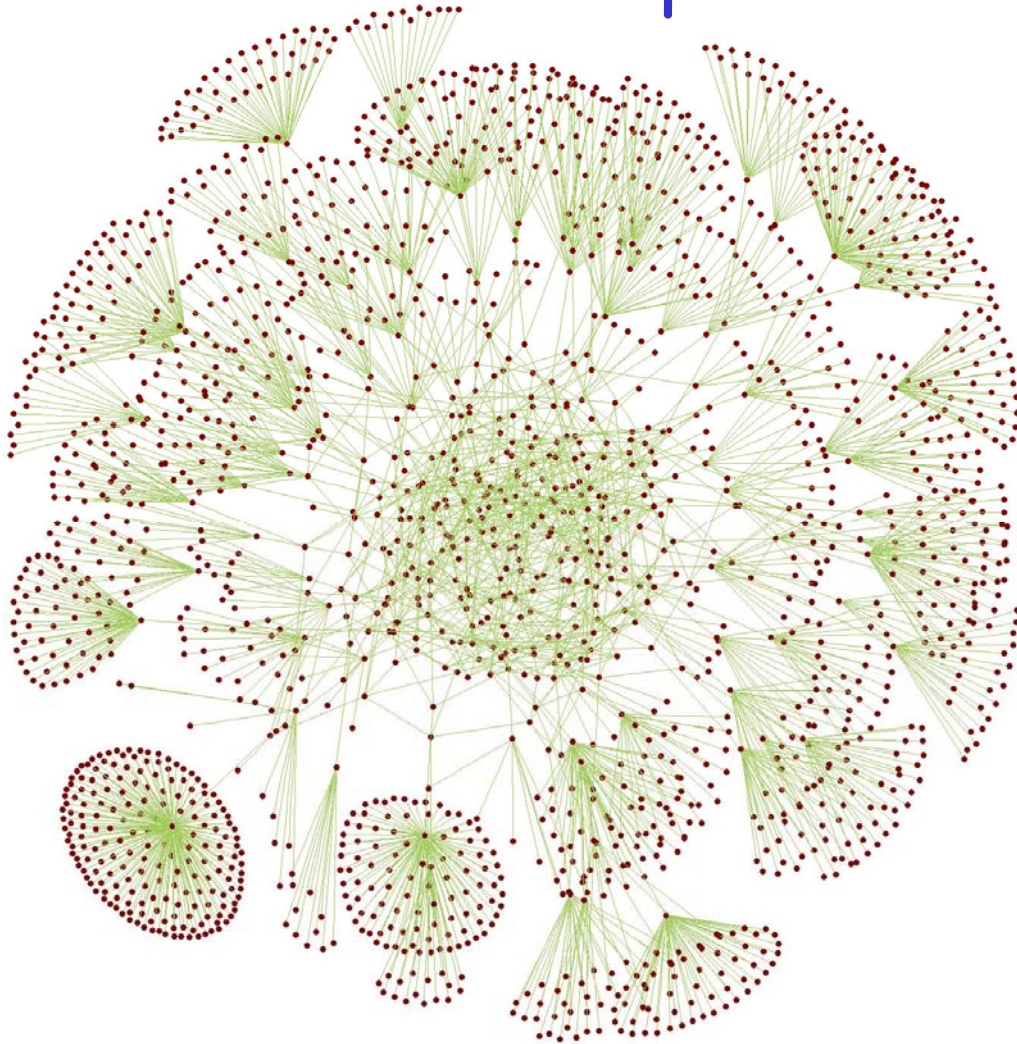


# CSCE 560

## Introduction to Computer Networking



Dr. Barry Mullins  
AFIT/ENG  
Bldg 642, Room 209  
255-3636 x7979

# Chapter 5: Outline

- ❑ 5.1 Introduction
- ❑ 5.2 Routing protocols
  - ❖ Link state
  - ❖ Distance vector
- ❑ 5.3 Intra-As routing in the Internet: OSPF
- ❑ 5.4 Routing among the ISPs: BGP
- ❑ 5.5 The SDN control plane
- ❑ 5.6 ICMP: The Internet Control Message Protocol
- ❑ 5.7 Network management and SNMP

# Network-layer Functions

Recall: two network-layer functions

---

- Forwarding:

- ❖ Move packets from router's input to appropriate router output

Data plane  
Chap 4

---

- Routing:

- ❖ Determine route taken by packets from source to destination

Control plane  
Chap 5

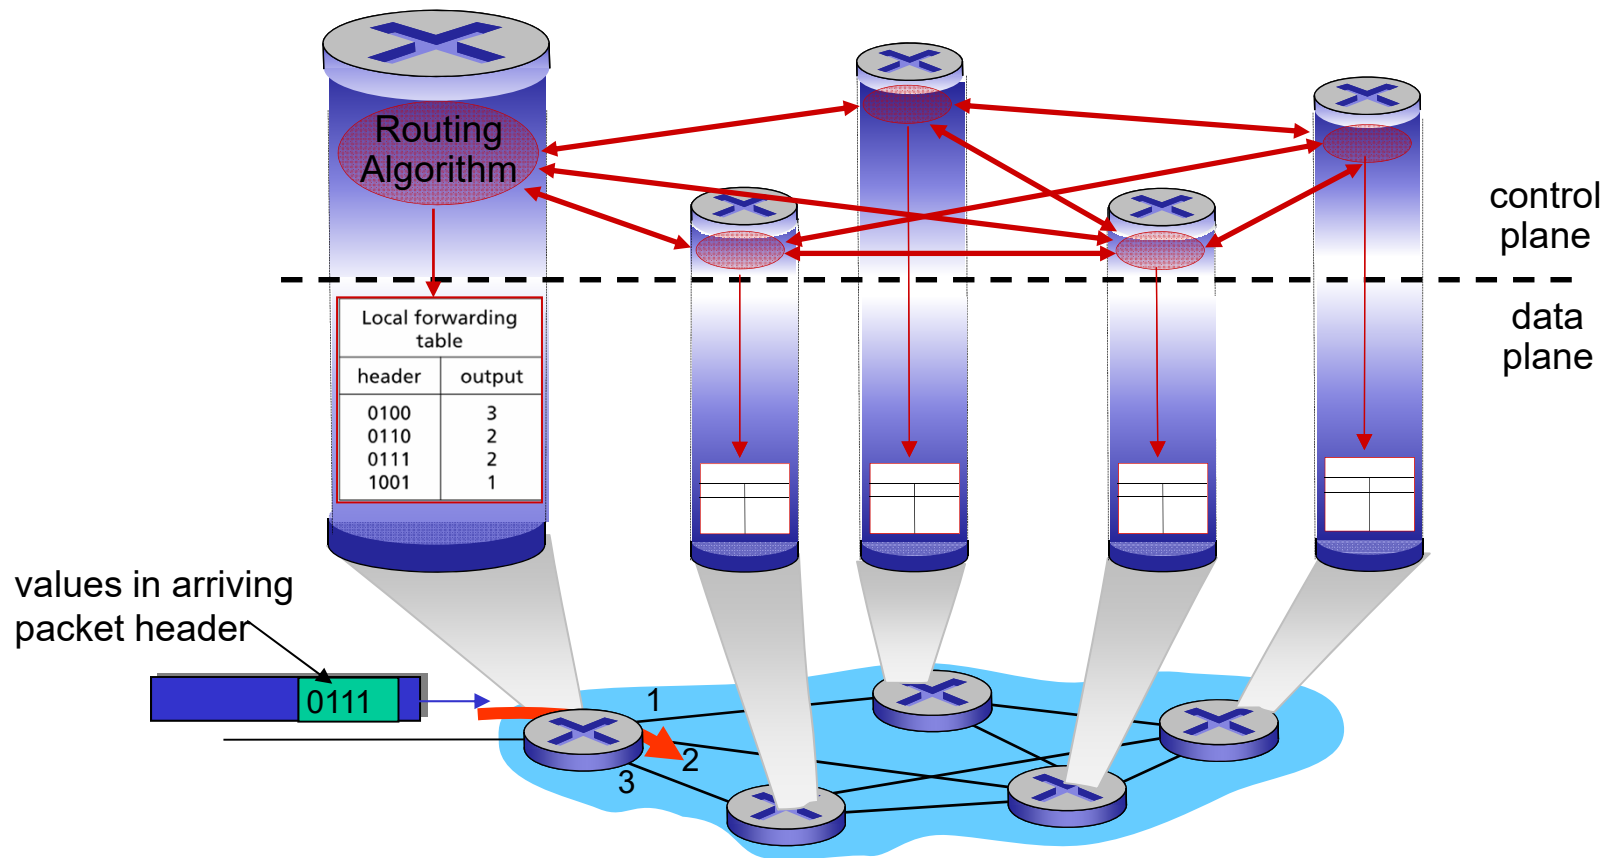
---

- Two approaches to structuring network control plane:

- ❖ **Traditional**: Per-router control
- ❖ **SDN**: Logically centralized control

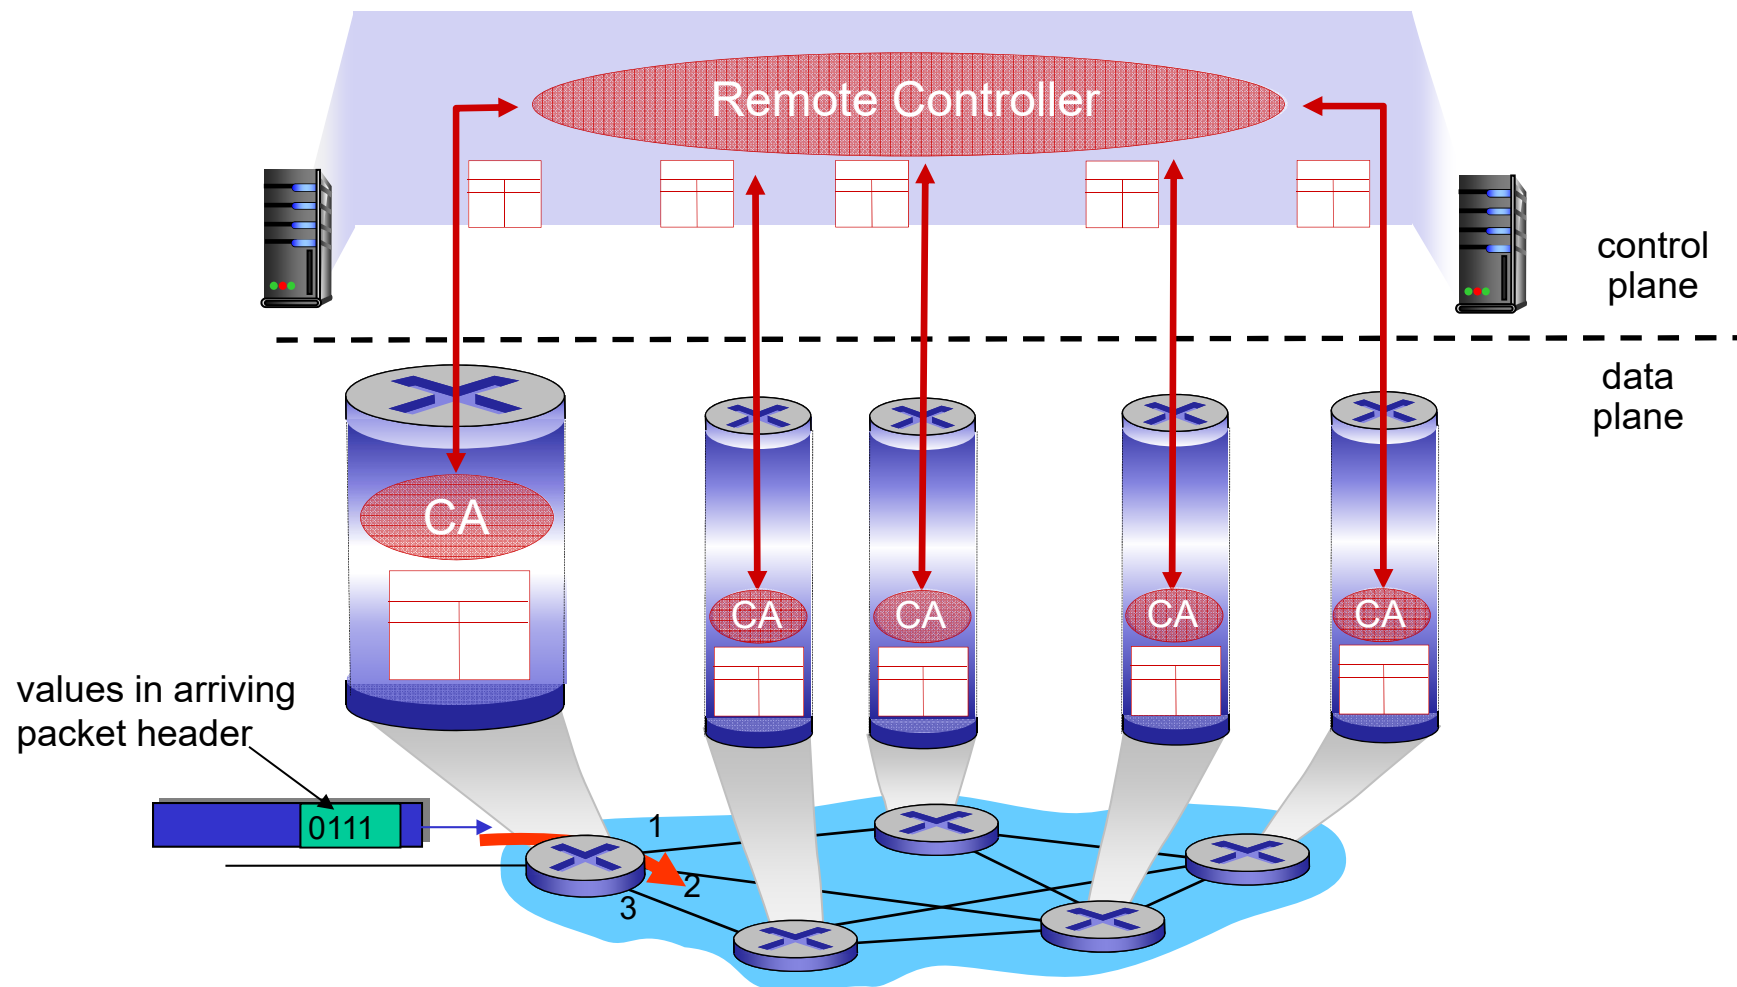
# Per-Router Control Plane

- Individual routing algorithm components in each and every router interact in the control plane to compute forwarding tables



# Logically Centralized Control Plane (SDN)

- A distinct (typically remote) controller interacts with local control agents (CAs)



# Chapter 5: Outline

- ❑ 5.1 Introduction
- ❑ 5.2 Routing protocols
  - ❖ Link state
  - ❖ Distance vector
- ❑ 5.3 Intra-As routing in the Internet: OSPF
- ❑ 5.4 Routing among the ISPs: BGP
- ❑ 5.5 The SDN control plane
- ❑ 5.6 ICMP: The Internet Control Message Protocol
- ❑ 5.7 Network management and SNMP

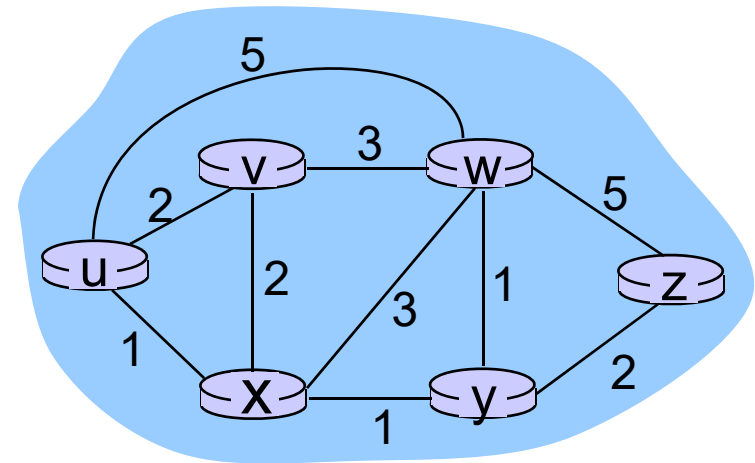
# Graph Abstraction

- Abstract network as a **graph**
  - ❖ Routers and end systems are nodes
  - ❖ Links and channels are edges

Graph:  $G = (N, E)$

$N$  = set of routers =  $\{ u, v, w, x, y, z \}$

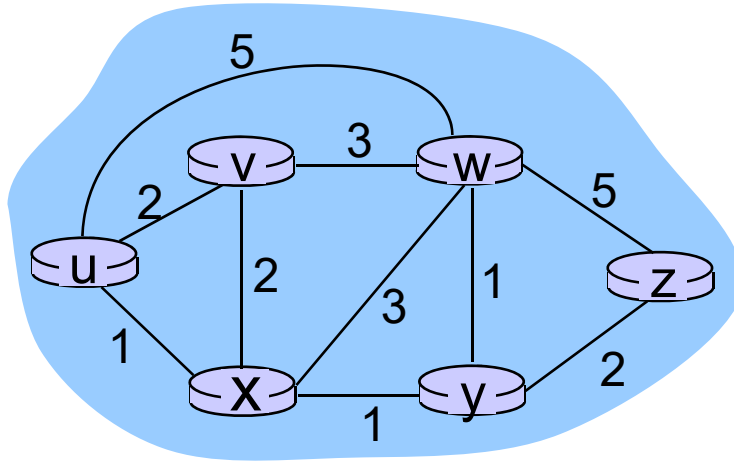
$E$  = set of links =  $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$



Remark: Graph abstraction is useful in other network contexts

Example: P2P, where  $N$  is set of peers and  $E$  is set of TCP connections

# Graph Abstraction: Costs



- $c(x,x')$  = cost of link  $(x,x')$ 
  - e.g.,  $c(w,z) = 5$
- Cost could always be 1, or inversely related to bandwidth

Cost of path  $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

Question: What's the least-cost path between u and z ?

Routing algorithm: algorithm that finds least-cost path



# Routing Algorithm Classification

Static or dynamic?

Static:

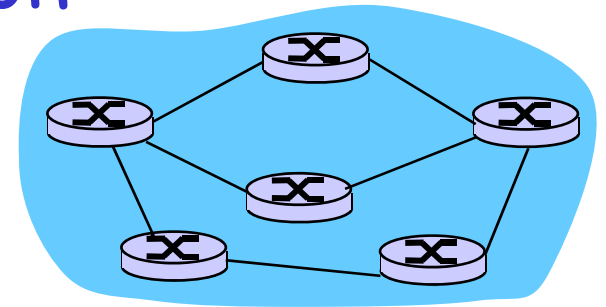
- ❑ Routes change very slowly over time
- ❑ A human updates the forwarding table 😊

Dynamic:

- ❑ Routes change more quickly
  - ❖ Periodic update
  - ❖ In response to link cost changes

# Routing Algorithm Classification

Global or decentralized information?



Global:

- ❑ All routers have complete global topology and link cost info
- ❑ Router has a map of the entire network before algorithm is run
- ❑ **"Link state" algorithms**
  - ❖ Flood (**broadcast**) routing (**link**) information to **all nodes**
  - ❖ Each router sends only the portion of the forwarding table that describes the state of its own links
  - ❖ Each router builds a picture of the entire network in its routing tables

# Chapter 5: Outline

- ❑ 5.1 Introduction
- ❑ 5.2 Routing protocols
  - ❖ Link state
  - ❖ Distance vector
- ❑ 5.3 Intra-As routing in the Internet: OSPF
- ❑ 5.4 Routing among the ISPs: BGP
- ❑ 5.5 The SDN control plane
- ❑ 5.6 ICMP: The Internet Control Message Protocol
- ❑ 5.7 Network management and SNMP

# A Link-State Routing Algorithm

## Dijkstra's algorithm

- Global net topology and link costs known to all nodes
  - ❖ All nodes have same info
  - ❖ Accomplished via "link state broadcast"
- Computes least cost paths from itself ("source") to all other nodes
  - ❖ Generates **forwarding table** for itself
- Iterative: after  $k$  iterations, know least cost path to  $k$  destinations

## Notation:

- $c(x,y)$ : link cost from node  $x$  to  $y$ 
  - ❖  $\infty$  if not direct neighbors
- $D(v)$ : current value of cost of path from source to dest.  $v$
- $p(v)$ : predecessor node along path from source to  $v$
- $N'$ : set of nodes whose least cost path definitively known
  - ❖ This set grows until all nodes are included

# Dijkstra's Algorithm for Source Node u

1 **Initialization:**

2  $N' = \{u\}$

3 for all nodes b

4 if b adjacent to u

5 then  $D(b) = c(u,b)$

6 else  $D(b) = \infty$

7

8 **Loop**

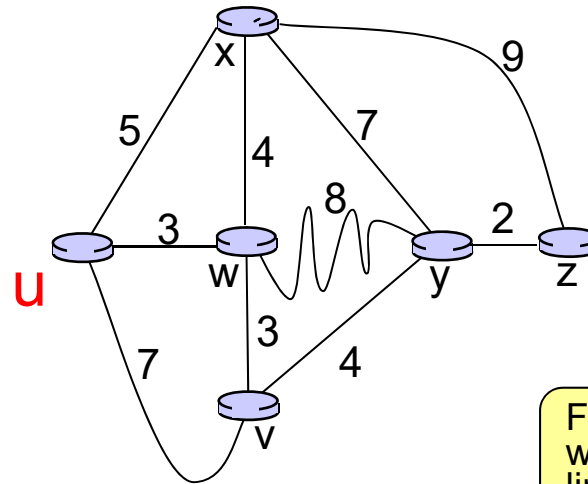
9 find e not in  $N'$  such that  $D(e)$  is a minimum

10 add e to  $N'$

11 update  $D(b)$  for all b adjacent to e and not in  $N'$  :

12  $D(b) = \min( D(b), D(e) + c(e,b) )$

13 **until all nodes in  $N'$**



Find neighbor  
with smallest  
link cost

New cost to b is either  
- old cost from u to b  
or  
- known shortest path cost from  
u to e plus cost from e to b

# Dijkstra's Algorithm: Example

Step	N'	D(v), p(v)	D(w), p(w)	D(x), p(x)	D(y), p(y)	D(z), p(z)
0	u	7,u	3,u	5,u	$\infty$	$\infty$
1	uw	6,w		5,u	11,w	$\infty$
2	uwx	6,w			11,w	14,x
3	uwxv			10,v		14,x
4	uwxvy				12,y	
5	uwxvyz					

## Loop

find e not in N' such that D(e) is a minimum  
add e to N'

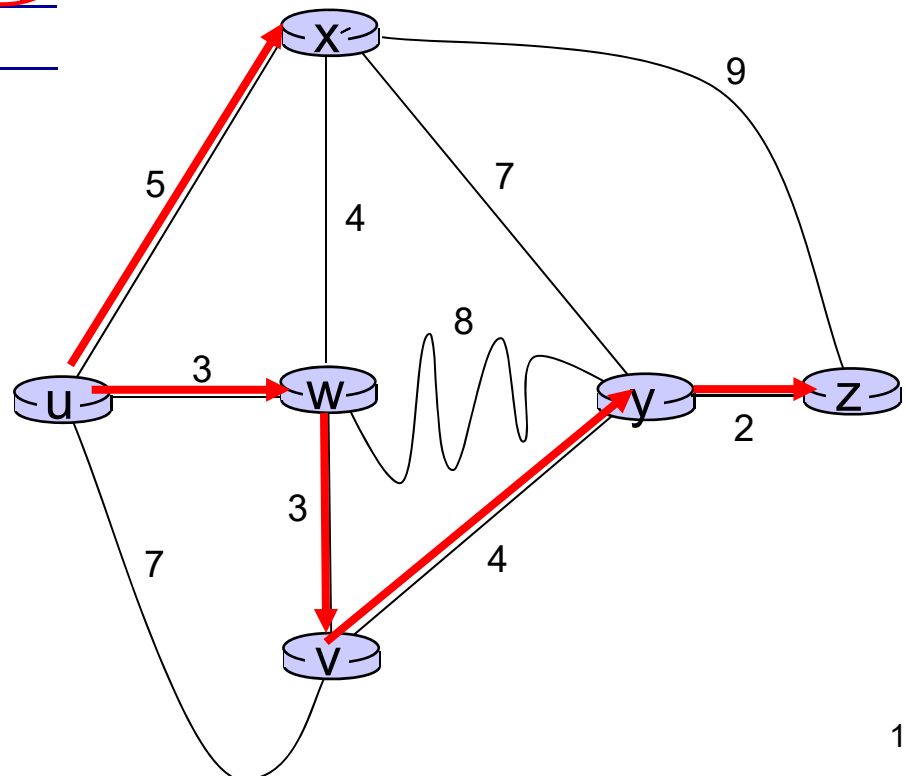
update D(b) for all b adjacent to e and not in N' :

$$D(b) = \min( D(b), D(e) + c(e,b) )$$

until all nodes in N'

## Notes:

- Construct shortest path tree by tracing predecessor nodes
- Ties can exist (can be broken arbitrarily)



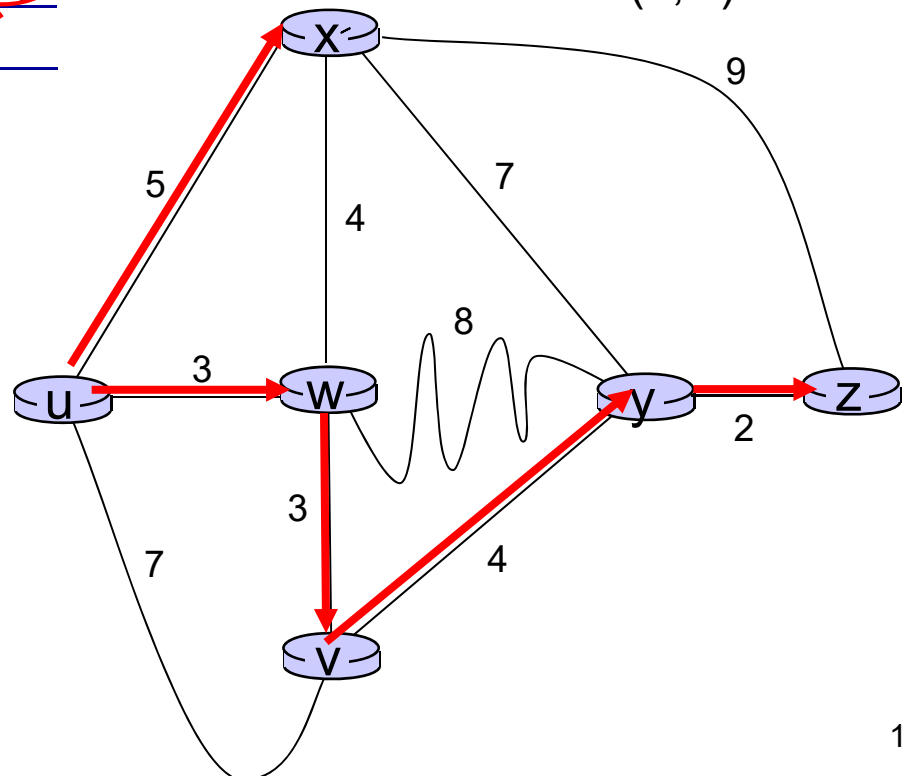
# Dijkstra's Algorithm: Example

Step	N'	D(v), p(v)	D(w), p(w)	D(x), p(x)	D(y), p(y)	D(z), p(z)
0	u	7,u	3,u	5,u	$\infty$	$\infty$
1	uw	6,w		5,u	11,w	$\infty$
2	uwx	6,w			11,w	14,x
3	uwxv				10,v	14,x
4	uwxvy					12,y
5	uwxvyz					

Follow the predecessor nodes to discover the appropriate links

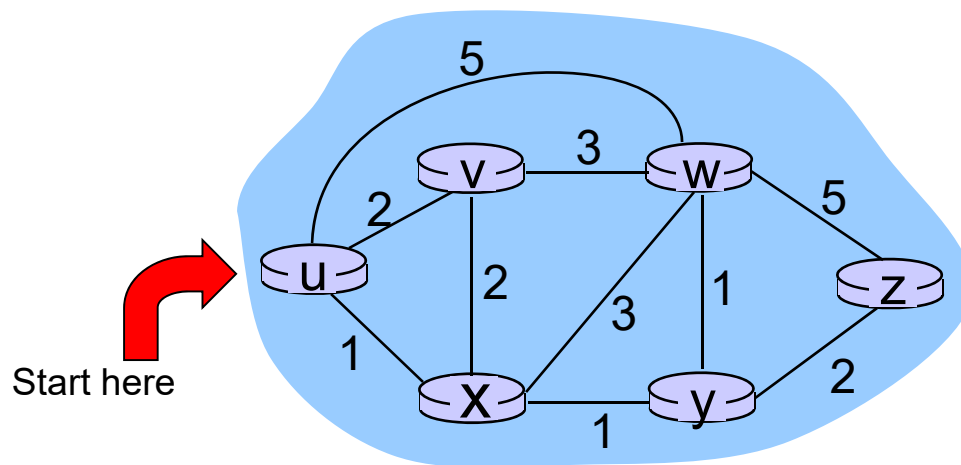
Resulting forwarding table in u:

Destination	Link
v	(u,w)
x	(u,x)
y	(u,w)
w	(u,w)
z	(u,w)



# Dijkstra's Algorithm: Another Example

Step	$N'$	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	u	2, u	5, u	1, u	$\infty$	$\infty$



- 1 **Initialization:**
- 2  $N' = \{u\}$
- 3 for all nodes  $b$
- 4 if  $b$  adjacent to  $u$
- 5 then  $D(b) = c(u, b)$
- 6 else  $D(b) = \infty$



# Dijkstra's Algorithm: Another Example

Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0	u	2,u	5,u	1,u	$\infty$	$\infty$
1	ux	2,u	4,x		2,x	$\infty$
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					

## Loop

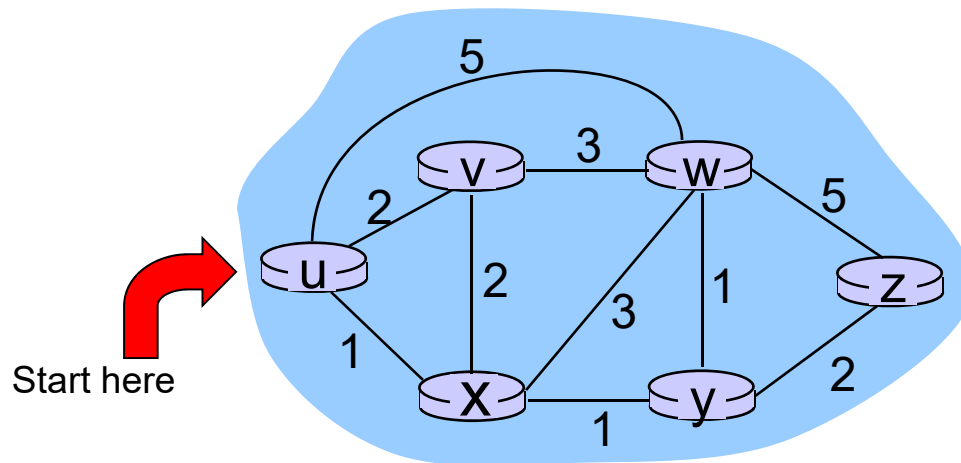
find e not in N' such that D(e) is a minimum

add e to N'

update D(b) for all b adjacent to e and not in N' :

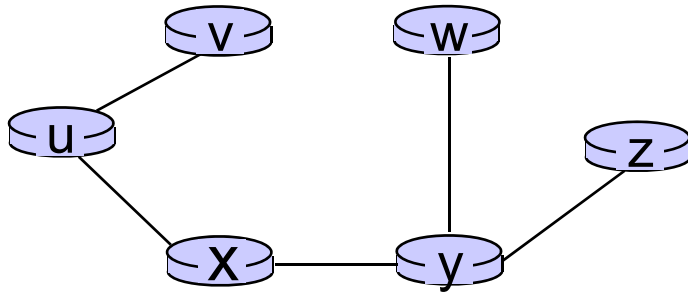
$$D(b) = \min( D(b), D(e) + c(e,b) )$$

**until all nodes in N'**



# Dijkstra's Algorithm: Resulting Table

Resulting shortest-path tree from u:



Resulting forwarding table in u:

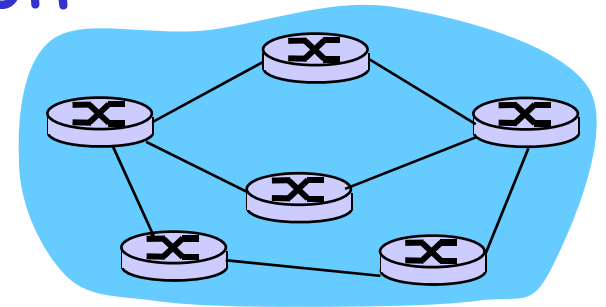
Destination	Link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
z	(u,x)

# Chapter 5: Outline

- ❑ 5.1 Introduction
- ❑ 5.2 Routing protocols
  - ❖ Link state
  - ❖ Distance vector
- ❑ 5.3 Intra-As routing in the Internet: OSPF
- ❑ 5.4 Routing among the ISPs: BGP
- ❑ 5.5 The SDN control plane
- ❑ 5.6 ICMP: The Internet Control Message Protocol
- ❑ 5.7 Network management and SNMP

# Routing Algorithm Classification

Global or decentralized information?



Decentralized:

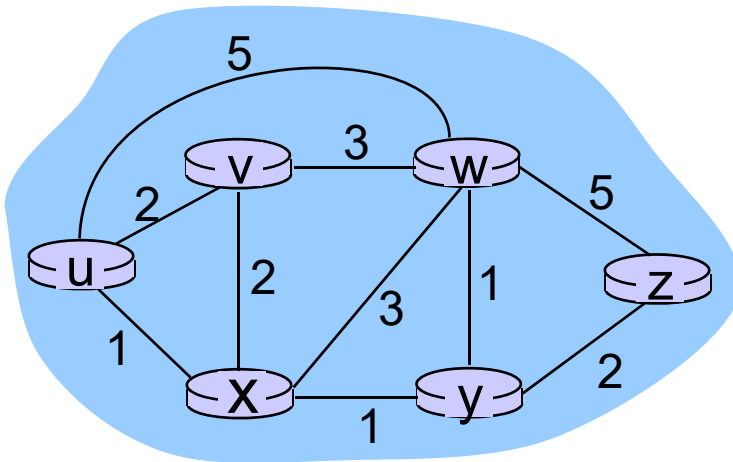
- ❑ Router knows physically-connected neighbors and link costs to neighbors
- ❑ Iterative, distributed process of computation by exchanging info with neighbors
- ❑ **"Distance vector" algorithms**
  - ❖ Each router sends all or some portion of its  
routing table only to its neighbors

# Bellman-Ford Equation

Define  $d_x(y) :=$  actual cost of least-cost **path** from  $x$  to  $y$

$$d_u(z) = \min_a \{ c(u,a) + d_a(z) \}$$

Known:  $d_v(z) = 5$ ,  $d_x(z) = 3$ ,  $d_w(z) = 3$



B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Node that achieves minimum is next  
hop in shortest path → add to forwarding table

# Distance Vector Algorithm

- $D_x(y)$  = estimate of least cost from  $x$  to  $y$
- Distance vector ( $D_x$ )
  - ❖ Cost estimates from  $x$  to all other nodes  $y$  in  $N$
- Node  $x$  maintains the following data
  1. Cost to each neighbor  $v$ :  $c(x,v)$
  2. Its distance vector (path costs to nodes that  $x$  knows about)
    - $D_x = [D_x(y): y \in N]$
  3. Its neighbors' distance vectors
    - $D_v = [D_v(y): y \in N]$

# Distance Vector Algorithm

## Basic idea:

- Each node periodically sends its own distance vector (DV) estimate to neighbors
- When a node  $x$  receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$

- Node  $x$  sends its DV to its neighbors if the DV changed
- Under natural conditions, the estimate  $D_x(y)$  converges to the actual least cost  $d_x(y)$

# Distance Vector Algorithm

## Iterative, asynchronous:

Each local iteration caused by:

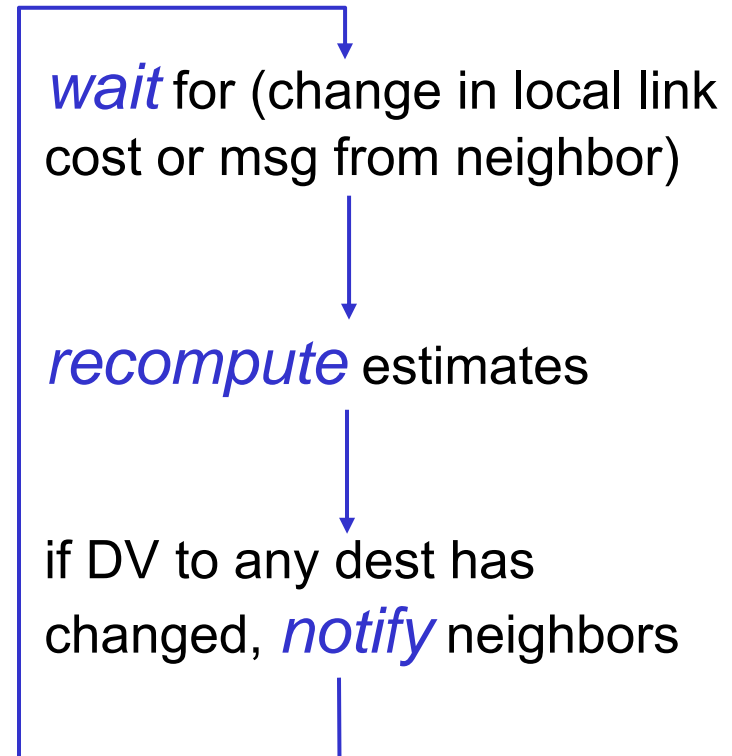
- ❖ Local link cost change
- ❖ DV update message from neighbor

## Distributed:

Each node notifies neighbors *only* when its DV changes

- ❖ Neighbors then notify their neighbors if necessary

## Each node:





# Distance Vector Algorithm

At all nodes, x:

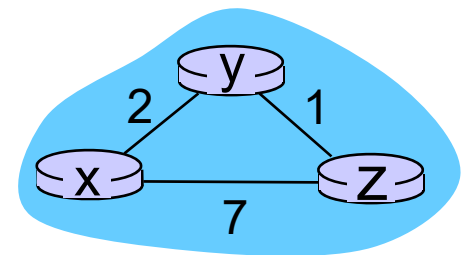
## 1 Initialization:

- 2 for all destinations y in N:
- 3  $D_x(y) = c(x,y)$  /\* if y is not a neighbor then  $c(x,y) = \text{infinity}$  \*/
- 4 for each neighbor w
- 5  $D_w(y) = \text{infinity}$  for all destinations y in N
- 6 for each neighbor w
- 7 send distance vector  $\mathbf{D}_x = [D_x(y): y \text{ in } N]$  to w

Distance vector →

	cost to		
	x	y	z
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞

from



# Distance Vector Algorithm

```
8 loop
9   wait (until I see a link cost change to neighbor w
10      or until I receive update from neighbor w)
11
12   For each y in N:
13      $D_x(y) = \min_v \{c(x,v) + D_v(y)\}$ 
14
15   If  $D_x(y)$  changed for any destination y
16     send distance vector  $D_x = [D_x(y): y \text{ in } N]$  to all neighbors
17 forever
```

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2 + 0, 7 + 1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

### node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	$\infty$	$\infty$	$\infty$
	z	$\infty$	$\infty$	$\infty$

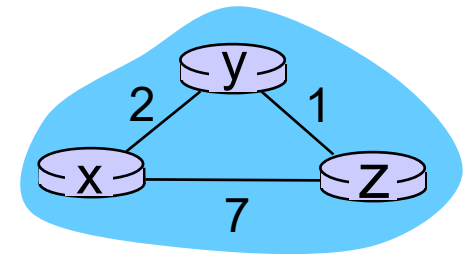
### node y table

		cost to		
		x	y	z
from	x	$\infty$	$\infty$	$\infty$
	y	2	0	1
	z	$\infty$	$\infty$	$\infty$

### node z table

		cost to		
		x	y	z
from	x	$\infty$	$\infty$	$\infty$
	y	$\infty$	$\infty$	$\infty$
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0



time

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2 + 0, 7 + 1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

### node x table

	cost to			
	x	y	z	
from x	0	2	7	
from y	$\infty$	$\infty$	$\infty$	
from z	$\infty$	$\infty$	$\infty$	

### node y table

	cost to			
	x	y	z	
from x	$\infty$	$\infty$	$\infty$	
from y	2	0	1	
from z	$\infty$	$\infty$	$\infty$	

### node z table

	cost to			
	x	y	z	
from x	$\infty$	$\infty$	$\infty$	
from y	$\infty$	$\infty$	$\infty$	
from z	7	1	0	

	cost to			
	x	y	z	
from x	0	2	3	
from y	2	0	1	
from z	7	1	0	

	cost to			
	x	y	z	
from x	0	2	7	
from y	2	0	1	
from z	7	1	0	

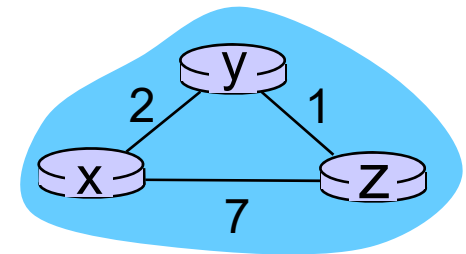
	cost to			
	x	y	z	
from x	0	2	7	
from y	2	0	1	
from z	3	1	0	

	cost to			
	x	y	z	
from x	0	2	3	
from y	2	0	1	
from z	3	1	0	

	cost to			
	x	y	z	
from x	0	2	3	
from y	2	0	1	
from z	3	1	0	

	cost to			
	x	y	z	
from x	0	2	3	
from y	2	0	1	
from z	3	1	0	

time



Resulting forwarding table in x:

Destination	Link
y	(x,y)
z	(x,y)

# Chapter 5: Outline

- ❑ 5.1 Introduction
- ❑ 5.2 Routing protocols
  - ❖ Link state
  - ❖ Distance vector
- ❑ 5.3 Intra-As routing in the Internet: OSPF
- ❑ 5.4 Routing among the ISPs: BGP
- ❑ 5.5 The SDN control plane
- ❑ 5.6 ICMP: The Internet Control Message Protocol
- ❑ 5.7 Network management and SNMP

# Making Routing Scalable

Our routing study thus far →  
idealization

- ❑ All routers identical
- ❑ Network “flat”  
... not true in practice

## Administrative autonomy

- ❑ Internet = network of networks
- ❑ Each network admin may want to control routing in own network

## Scale:

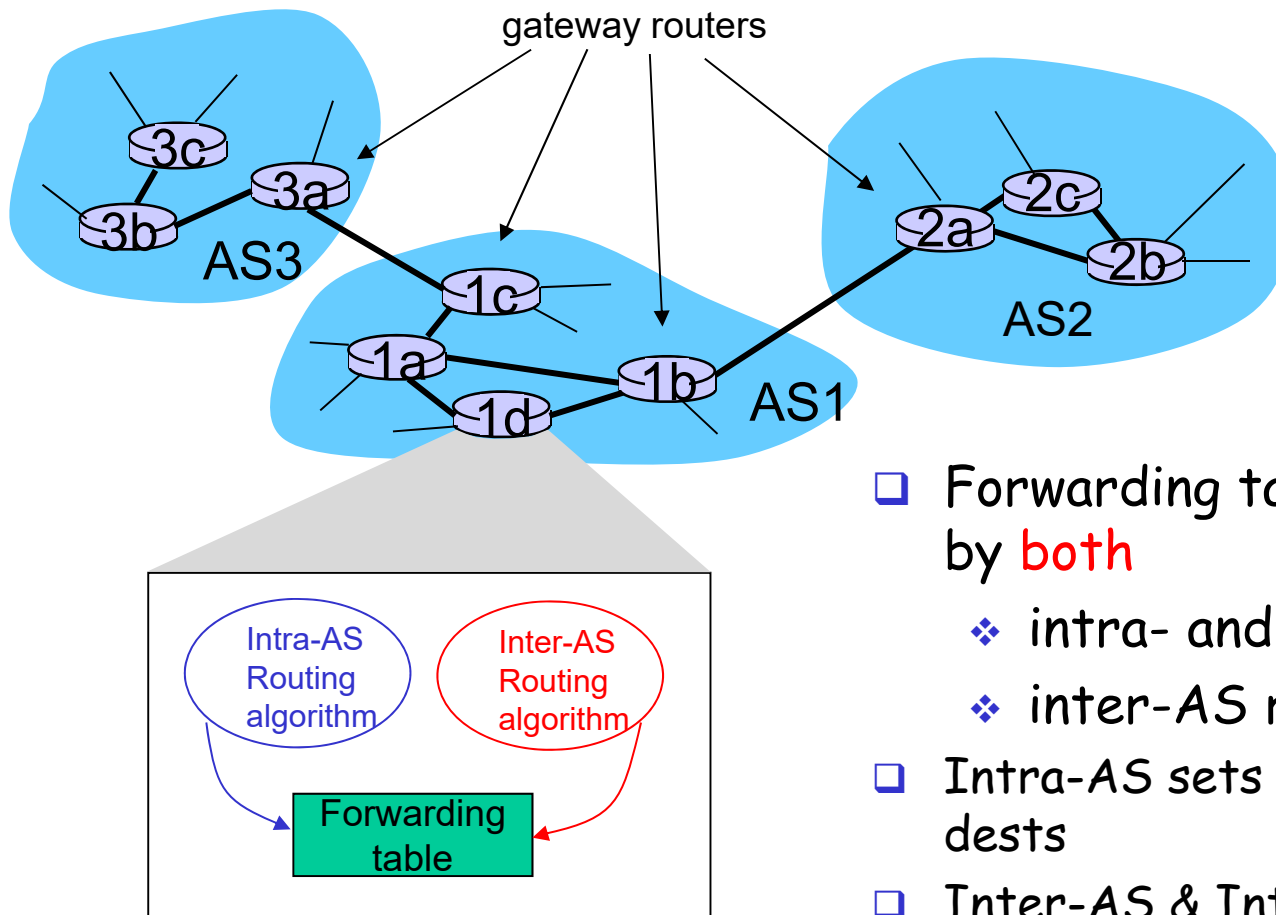
With billions of destinations:

- ❑ Can't store all dest's in routing tables!
- ❑ Routing table exchange would swamp links!

# Internet Approach to Scalable Routing

- Aggregate routers into regions called "autonomous systems" (AS) (aka domains)
  - ❖ Each AS is assigned an AS Number (ASN) by IANA
- **Intra-AS routing**
  - ❖ Routers in same AS run same routing protocol
    - "INTRA-AS" routing protocol
  - ❖ Routers in different AS can run different intra-AS routing protocol
  - ❖ An intra-AS router only needs to know about the other routers in the AS
- **Inter-As routing**
  - ❖ Routing among AS's
  - ❖ Gateways perform inter-domain routing (as well as intra-domain routing)
- Gateway router
  - ❖ At "edge" of its own AS
  - ❖ Direct link to router in another AS
- AFIT ASN = AS133
  - ❖ Not currently in use
- WPAFB ASN = AS132

# Interconnected ASes



- ❑ Forwarding table is configured by **both**
  - ❖ intra- and
  - ❖ inter-AS routing algorithms
- ❑ Intra-AS sets entries for internal dests
- ❑ Inter-AS & Intra-AS sets entries for external dests



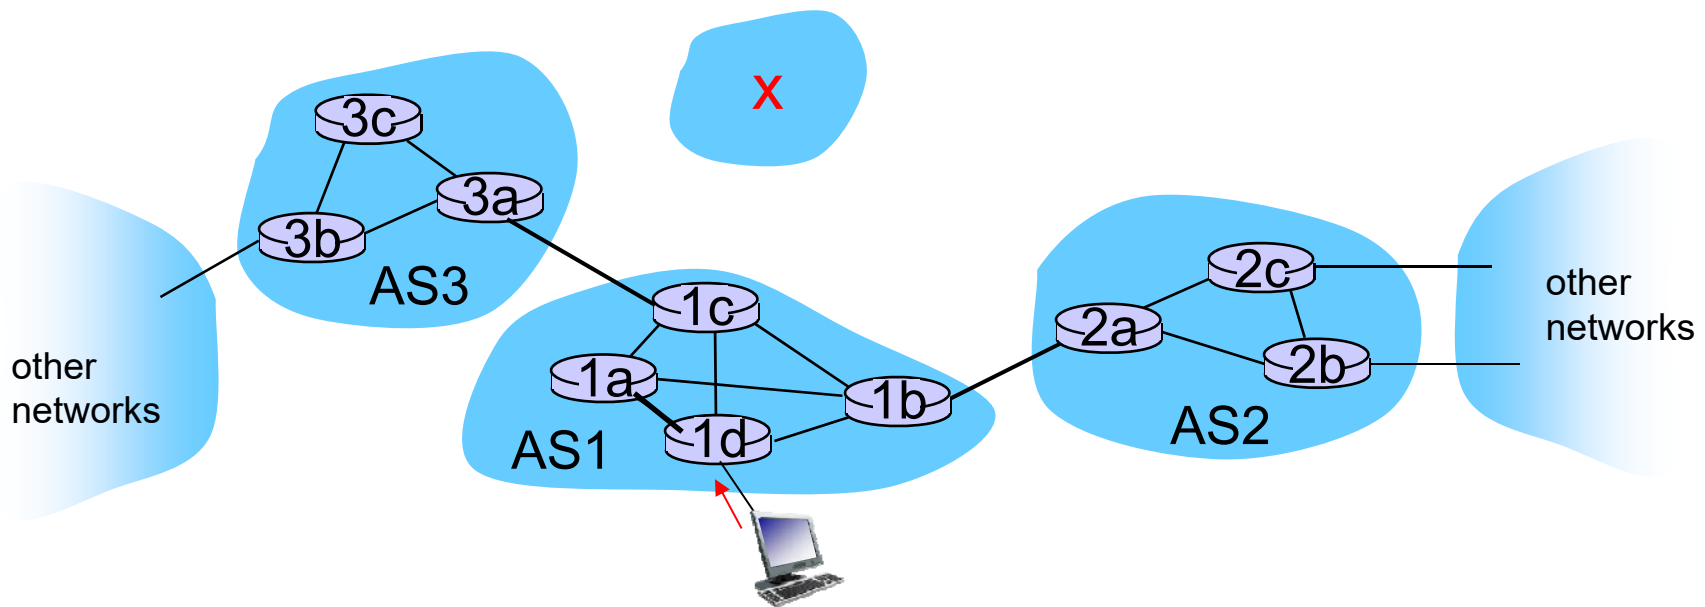
# Inter-AS tasks

- Suppose router 1d in AS1 receives datagram for a dest outside of AS1 → subnet X
  - ❖ 1d should forward packet towards one of the gateway routers, but which one?

## AS1 needs:

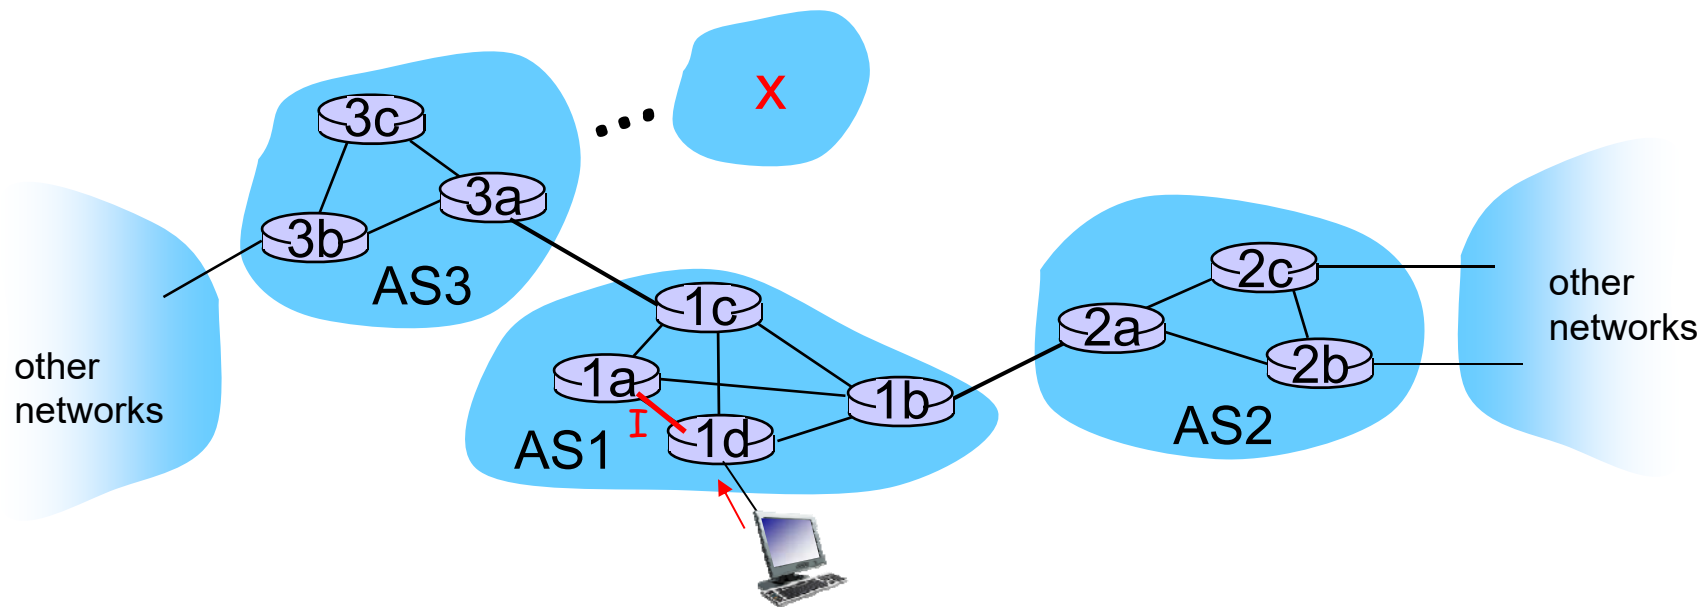
1. To learn which dests are reachable through AS2 and AS3
2. To propagate this reachability info to all routers in AS1

Job of inter-AS routing!



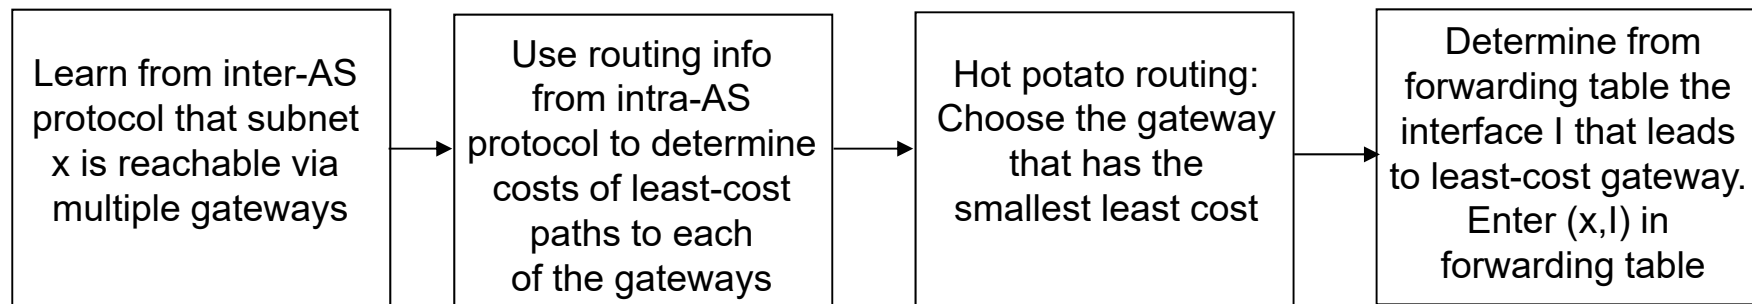
# Example: Setting Forwarding Table in Router 1d

- Suppose AS1 learns from the inter-AS protocol that subnet **x** is reachable through AS3 (gateway 1c) but not from AS2
- Inter-AS protocol propagates reachability info to internal routers
- Router 1d determines from intra-AS routing info that its interface **I** is on the least cost path to 1c
- Puts in forwarding table entry **(x,I)**



# Example: Choosing Among Multiple ASes

- ❑ Now suppose AS1 learns from the inter-AS protocol that subnet **x** is reachable from AS3 **and** from AS2
- ❑ To configure forwarding table, router 1d must determine towards which gateway it should forward packets for dest **x**
- ❑ This is also the job of intra-AS routing protocol
- ❑ **Hot potato routing:**
  - ❖ Send packet towards closest of two gateway routers



# Intra-AS Routing

- Also known as **Interior Gateway Protocols (IGP)**
- Most common Intra-AS routing protocols:
  - ❖ OSPF: Open Shortest Path First
    - Dijkstra (Link State)
  - ❖ RIP: Routing Information Protocol
    - Bellman-Ford (Distance Vector)
  - ❖ IGRP: Interior Gateway Routing Protocol (Cisco proprietary)

# OSPF (Open Shortest Path First)

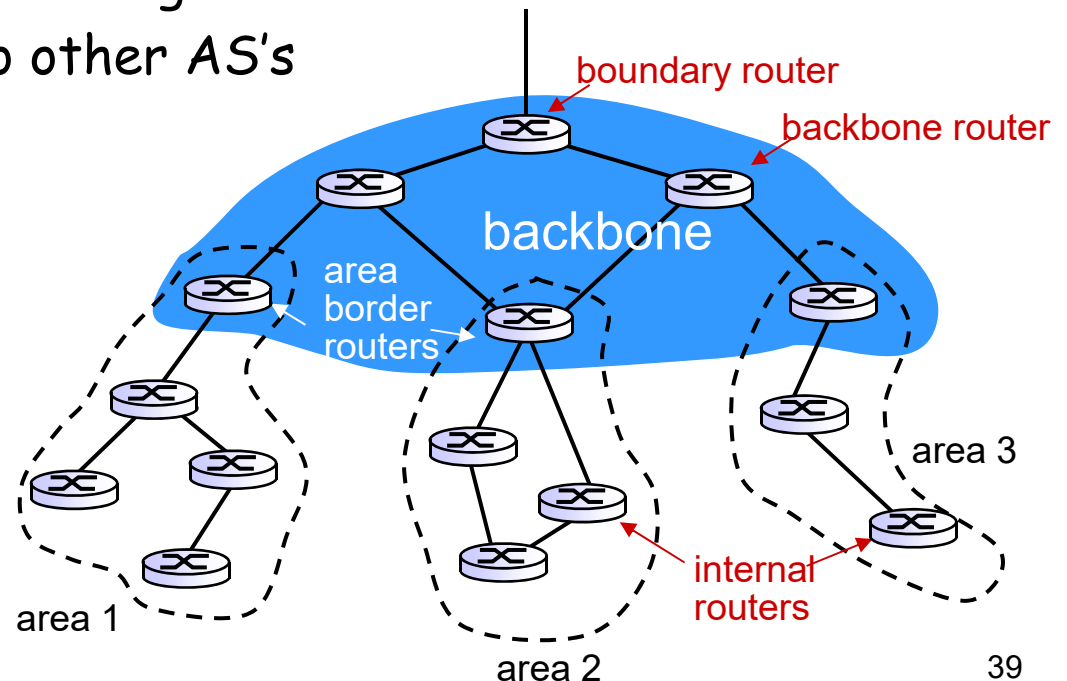
- ❑ "Open": publicly available
- ❑ Uses Link State algorithm
- ❑ OSPF advertisement carries one entry (link state) per neighbor router
- ❑ Advertisements disseminated to entire AS or area (via flooding)
  - ❖ Carried in OSPF messages directly over IP (upper layer = 89)
  - ❖ Sent at least every 30 minutes
  - ❖ Routers detect each other using HELLO OSPF protocol packet broadcasts
  - ❖ Routers then select a designated router (DR) which acts as a hub to reduce traffic between routers
    - DR maintains a complete topology table of the network and sends the updates to the other routers via multicast
    - Every time a router sends an update, it sends it to the DR
      - DR then sends the update out to all other routers in the area

# OSPF "Advanced" Features

- ❑ Link costs set by administrator instead of all links = 1 (as in RIP)
- ❑ **Security**: all OSPF messages authenticated (to prevent malicious intrusion)
- ❑ **Multiple** same-cost **paths** allowed (only one path in RIP)
- ❑ Support for **Hierarchical** OSPF in large domains
- ❑ Typically deployed in upper-tier ISPs

# Hierarchical OSPF

- ❑ **Two-level hierarchy:** 1. local area, 2. backbone
  - ❖ Link-state advertisements only in area
  - ❖ Each node has detailed area topology; only know direction (shortest path) to nets in other areas
- ❑ **Area border routers:** "summarize" distances to nets in own area, advertise to other area border routers
- ❑ **Backbone routers:** run OSPF routing limited to backbone
- ❑ **Boundary routers:** connect to other AS's



# Chapter 5: Outline

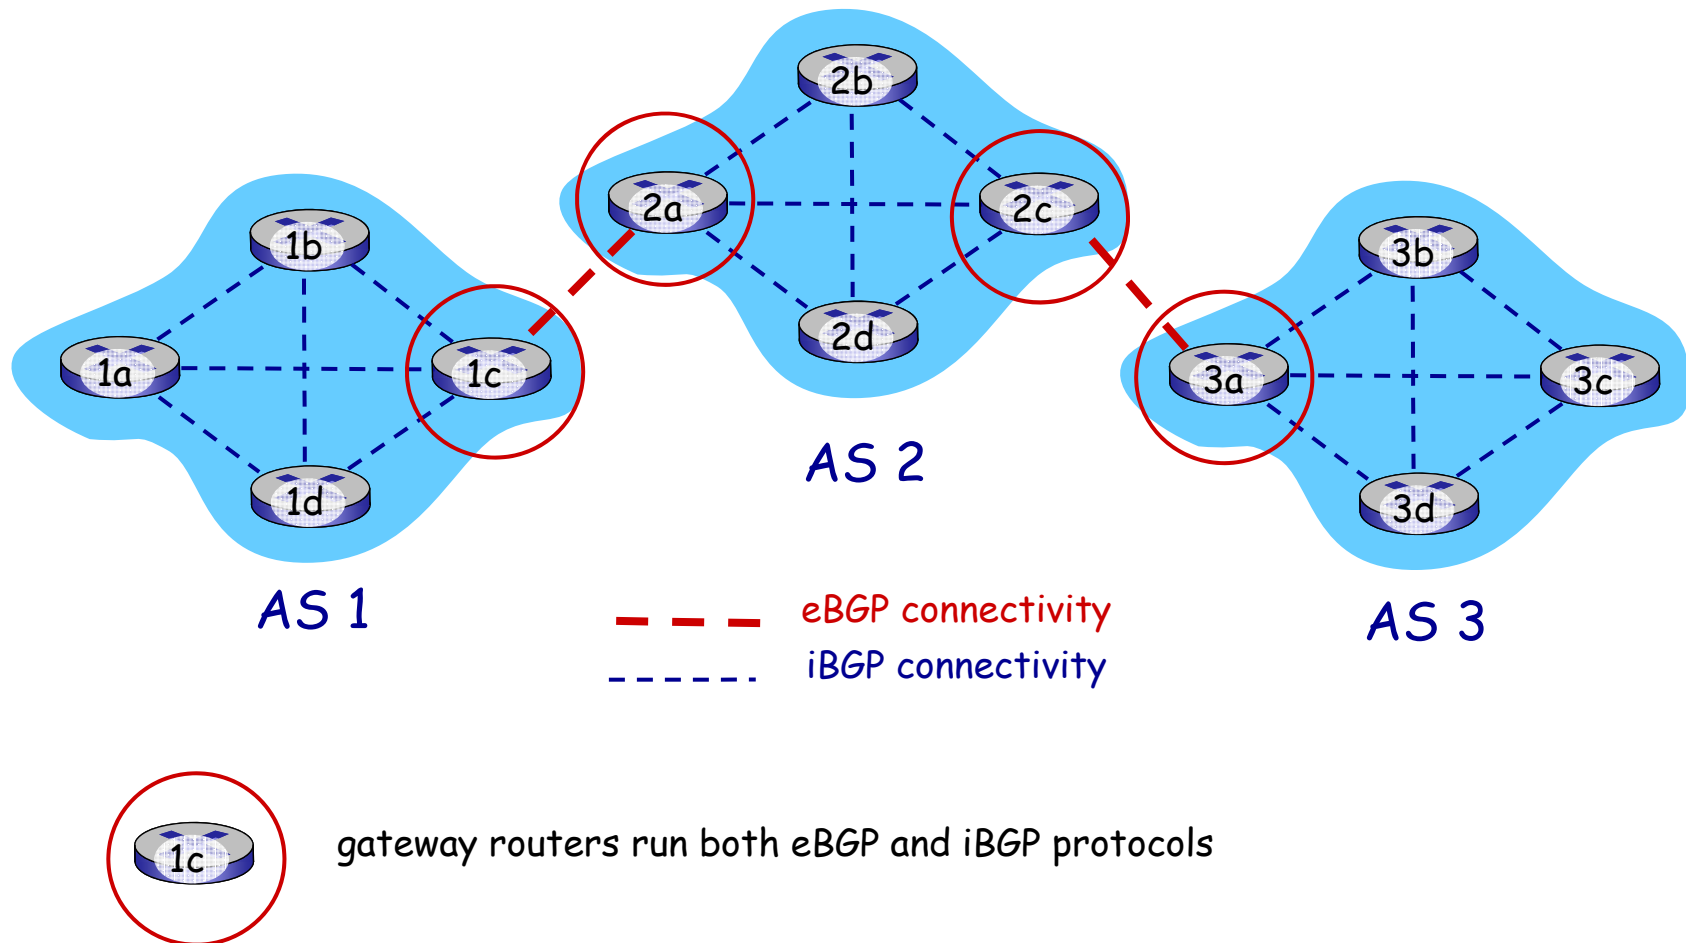
- ❑ 5.1 Introduction
- ❑ 5.2 Routing protocols
  - ❖ Link state
  - ❖ Distance vector
- ❑ 5.3 Intra-As routing in the Internet: OSPF
- ❑ 5.4 Routing among the ISPs: BGP
- ❑ 5.5 The SDN control plane
- ❑ 5.6 ICMP: The Internet Control Message Protocol
- ❑ 5.7 Network management and SNMP



# Internet INTER-AS Routing: BGP

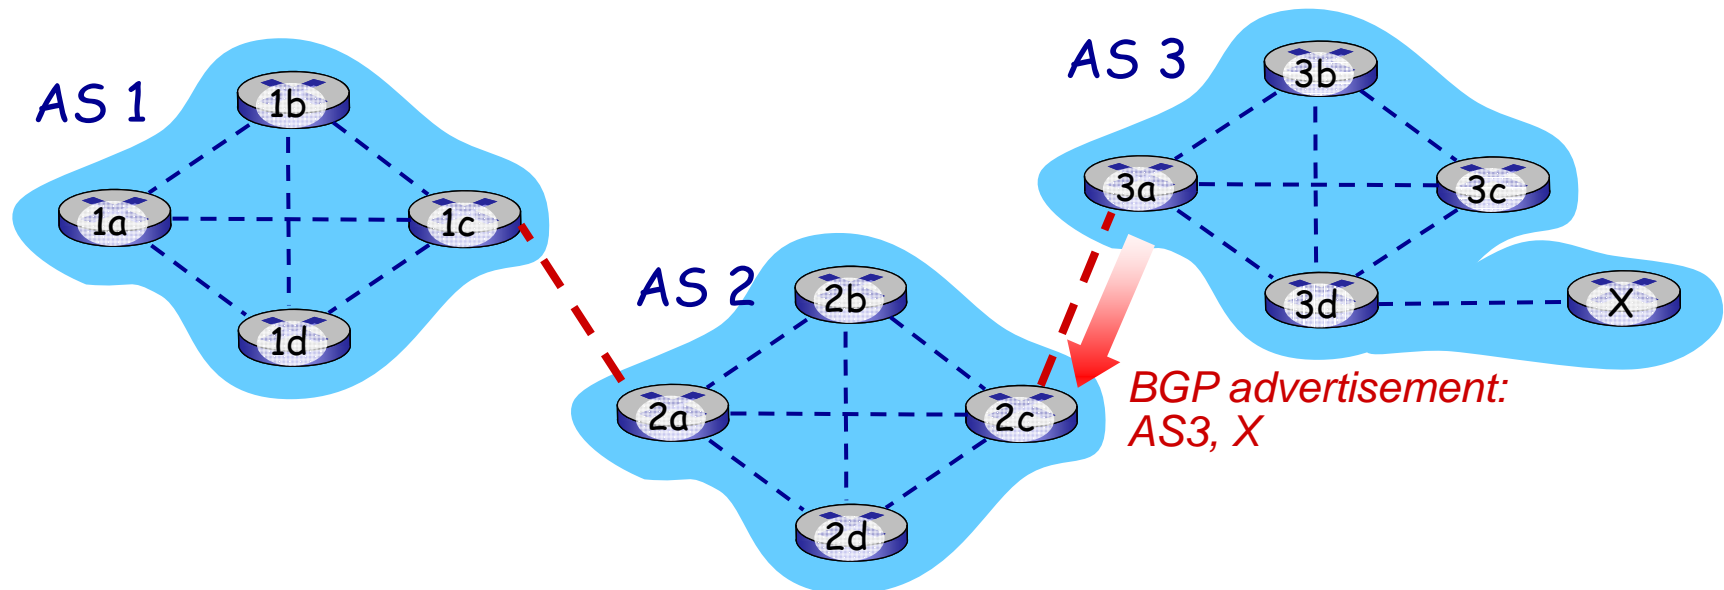
- **BGP (Border Gateway Protocol):** the de facto standard
- BGP provides each AS a means to:
  1. **eBGP:** Obtain subnet reachability info from neighboring (**external**) ASs
  2. **iBGP:** Propagate reachability info to all routers **internal** to the AS
  3. Determine “good” routes to subnets based on reachability information and policy
- Allows a subnet to advertise its existence to rest of the Internet:  
*“I am here”*

# eBGP, iBGP Connections



# BGP Basics

- **BGP session**: Pairs of routers (BGP peers) exchange routing info over semi-permanent TCP connections (port 179)
  - ❖ BGP sessions do not correspond to physical links
  - ❖ Keep-alive msgs sent every 30 seconds to maintain connection
- When AS3 gateway router 3a advertises path **AS3,X** to AS2 gateway router 2c:
  - ❖ AS3 **promises** to AS2 it will forward datagrams towards X



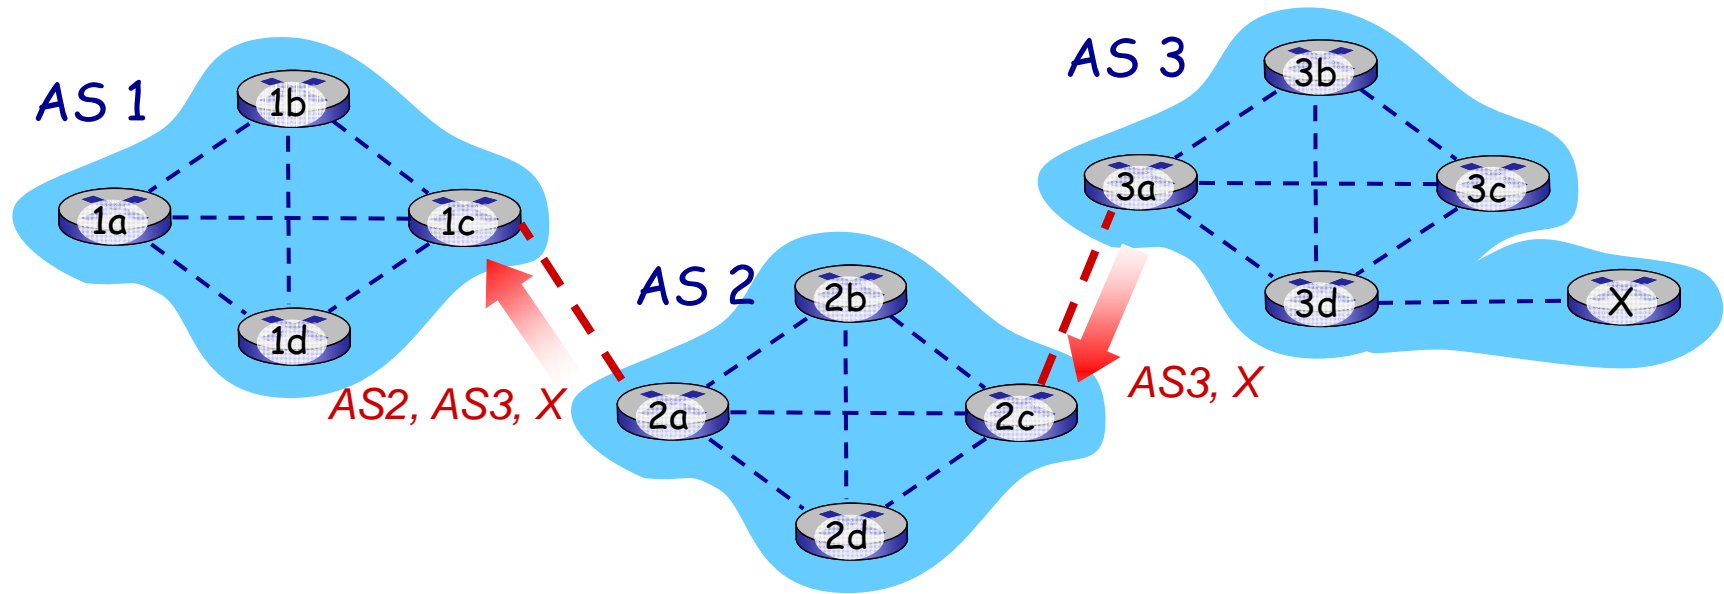
# Path Attributes and BGP Routes

- ❑ Advertised prefix includes BGP attributes
  - ❖ prefix + attributes = "route"
- ❑ Three mandatory "well known" attributes:
  - ❖ **Origin** - The origin of a BGP route
    - "i" (code 0) - internal - route aggregation (network command used)
    - "e" (code 1) - external - obsolete
    - "?" (code 2) - origin is unknown or route was redistributed/aggregated/incomplete
  - ❖ **AS-PATH**: contains the ASs through which the advertisement for the prefix passed: AS67 AS17
    - Useful for detecting loops
  - ❖ **NEXT-HOP**: IP address of the next-hop router
    - Sending router sets the next-hop field to its own IP address
    - Used to
      - establish BGP TCP sessions
      - determine least-cost path to closest link by internal routing algorithms

# Path Attributes and BGP Routes

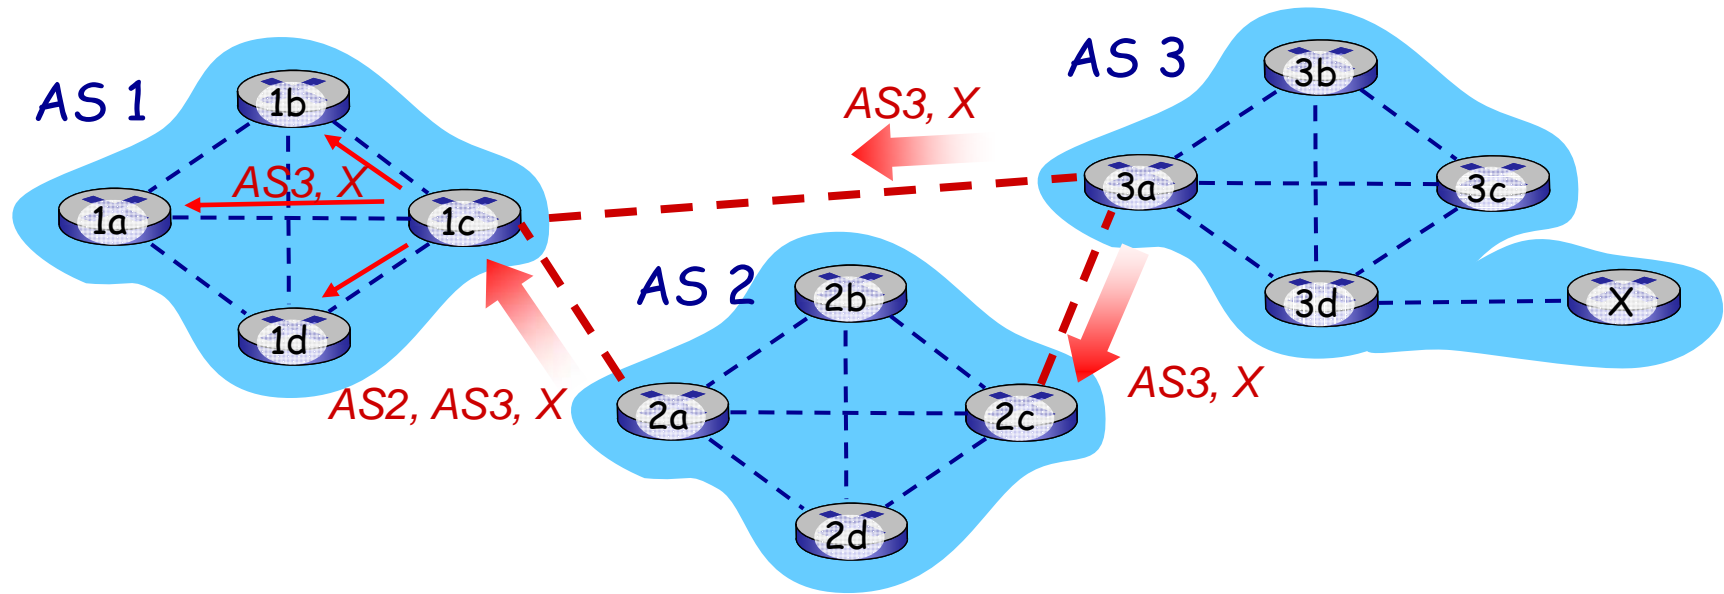
- Gateway router receiving route advertisement uses **import policy** to accept/decline
  - ❖ Why decline?
    - The AS may not want to send traffic through one of the ASs listed in AS-PATH or
    - It detected a loop
  
- **Policy-based** routing

# BGP Path Advertisement



- ❑ AS2 router 2c receives path advertisement **AS3,X** (via eBGP) from AS3 router 3a
- ❑ Based on AS2 policy, AS2 router 2c accepts path AS3,X, propagates (via iBGP) to all AS2 routers
- ❑ Based on AS2 policy, AS2 router 2a advertises (via eBGP) path **AS2, AS3, X** to AS1 router 1c

# BGP Path Advertisement



- Gateway router (e.g., 1c) may learn about **multiple** paths to destination
  - ❖ **AS2,AS3,X** from 2a
  - ❖ **AS3,X** from 3a
  - ❖ Based on policy, AS1 gateway router 1c chooses path AS3,X, and advertises path within AS1 via iBGP

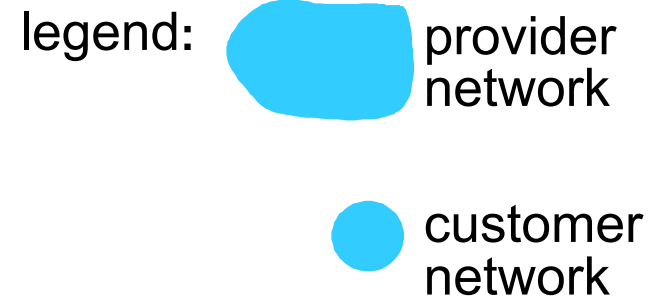
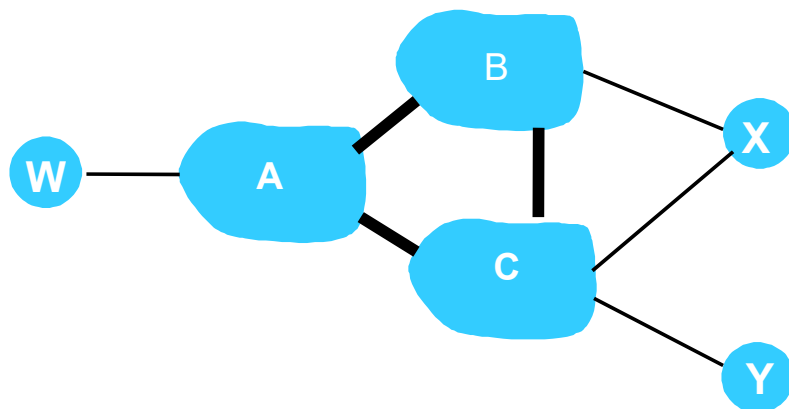
# BGP Route Selection

- Router may learn about more than 1 route to destination AS, selects route based on:
  1. Local preference value attribute: **policy decision**
  2. Shortest AS-PATH
  3. Closest NEXT-HOP router: hot potato routing
  4. "Additional criteria"



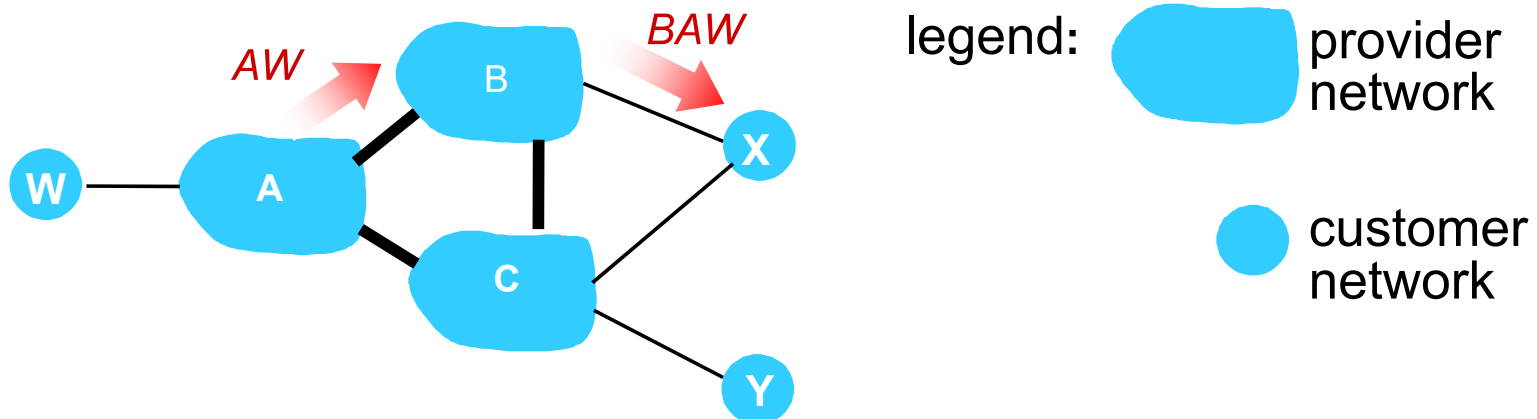
# BGP Routing Policy

- A,B,C are provider networks (ASs)
- W,X,Y are customers (of provider networks) and also ASs
- X is dual-homed: attached to two networks
  - ❖ X does not want to route from B to C
  - ❖ .. so X will not advertise to B a route to C



# BGP Routing Policy

- ❑ A advertises to B the path *AW*
- ❑ B advertises to X the path *BAW*
- ❑ Should B advertise to C the path *BAW*?
  - ❖ No way! B gets no "revenue" for routing *CBAW* since neither *W* nor *C* are B's customers
  - ❖ B wants to force C to route to *w* via *A*
  - ❖ B wants to route only to/from its customers!
- ❑ Rule of thumb - if traffic is for ISP's customers → forward



# Why Different Intra-, Inter-AS Routing?

## Policy:

- ❑ Inter-AS: admin wants control
  - ❖ over how its traffic routed and
  - ❖ who routes through its net
- ❑ Intra-AS: single admin
  - ❖ no policy decisions needed

## Scale:

- ❑ Hierarchical routing saves table size and reduces update traffic

## Performance:

- ❑ Intra-AS: can focus on **performance**
- ❑ Inter-AS: **policy** may dominate over performance

# Chapter 5: Outline

- ❑ 5.1 Introduction
- ❑ 5.2 Routing protocols
  - ❖ Link state
  - ❖ Distance vector
- ❑ 5.3 Intra-As routing in the Internet: OSPF
- ❑ 5.4 Routing among the ISPs: BGP
- ❑ 5.5 The SDN control plane
- ❑ 5.6 ICMP: The Internet Control Message Protocol
- ❑ 5.7 Network management and SNMP

# Historical Approach To Routing

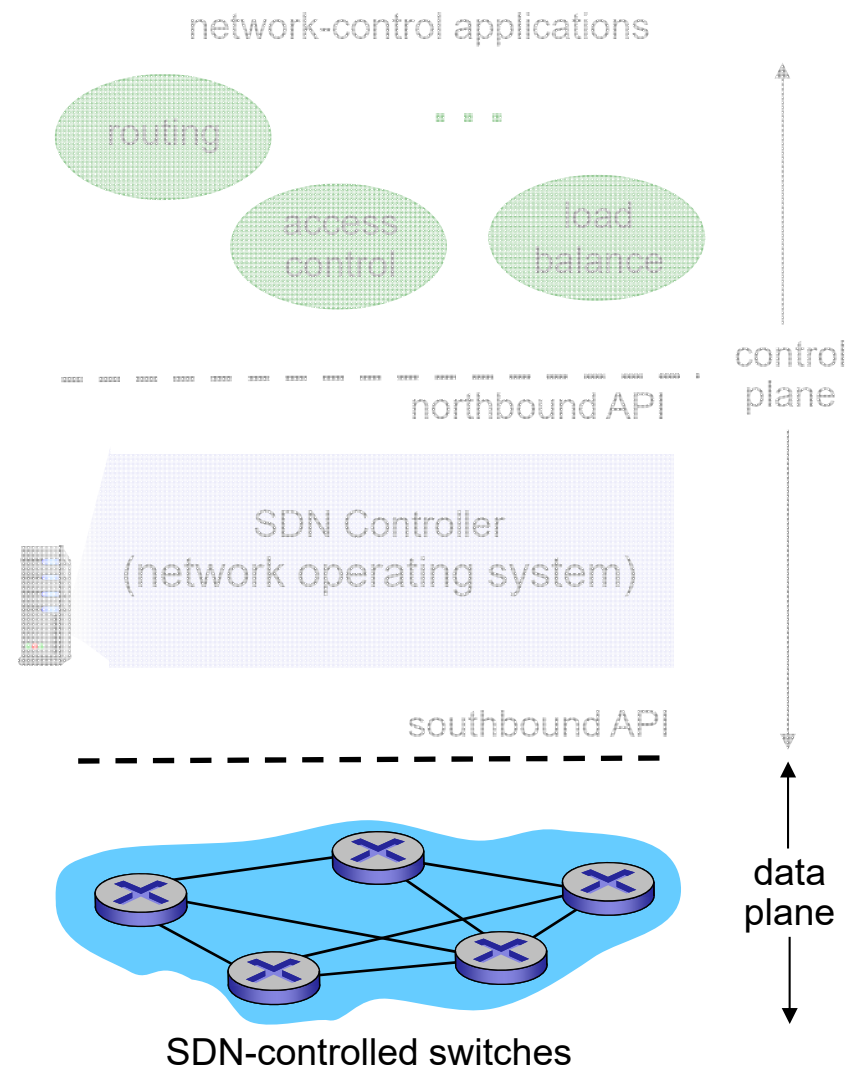
- ❑ Internet network layer: historically has been implemented via distributed, per-router approach
- ❑ Monolithic router contains
  - ❖ Switching hardware,
  - ❖ Runs proprietary implementation of Internet standard protocols (RIP, OSPF, BGP) in **proprietary** router OS (e.g., Cisco IOS)
- ❑ Different “middleboxes” for different network layer functions
  - ❖ Firewalls, routers, switches, load balancers, NAT boxes, ...

# Why a Logically Centralized Control Plane in SDN?

- ❑ Easier network management
  - ❖ Avoid router misconfigurations
  - ❖ Greater flexibility of traffic flows
- ❑ Table-based forwarding allows “programming” routers
  - ❖ Centralized “programming” easier
    - Compute tables centrally and distribute
  - ❖ Distributed (historical) “programming” more difficult
    - Compute tables as result of distributed algorithm (protocol) implemented in each and every router
- ❑ Open (non-proprietary) implementation of control plane

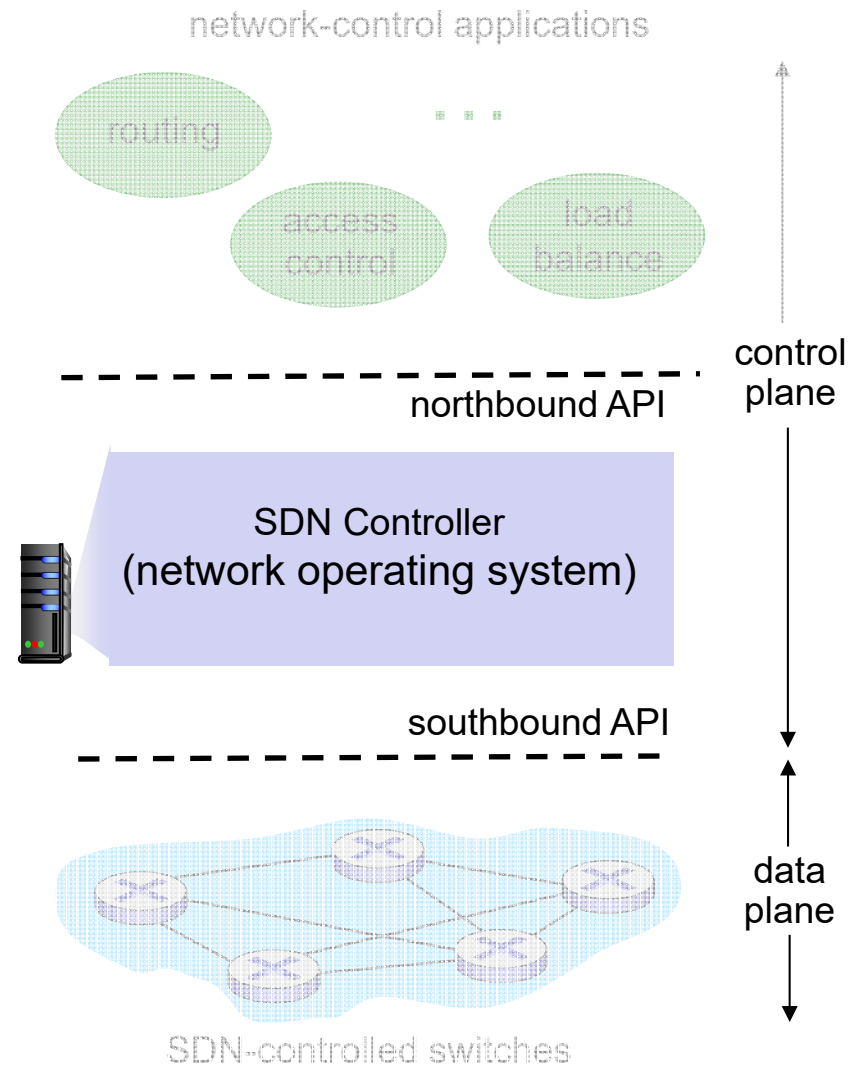
# SDN Perspective: Data Plane Switches

- ❑ Fast, simple, commodity switches implementing generalized data-plane forwarding in hardware
- ❑ Switch flow table computed and installed by controller
- ❑ API for table-based switch control (e.g., OpenFlow)
  - ❖ Defines what is controllable and what is not
- ❑ Protocol for communicating with controller
  - ❖ OpenFlow
    - TCP port 6653



# SDN Perspective: SDN Controller

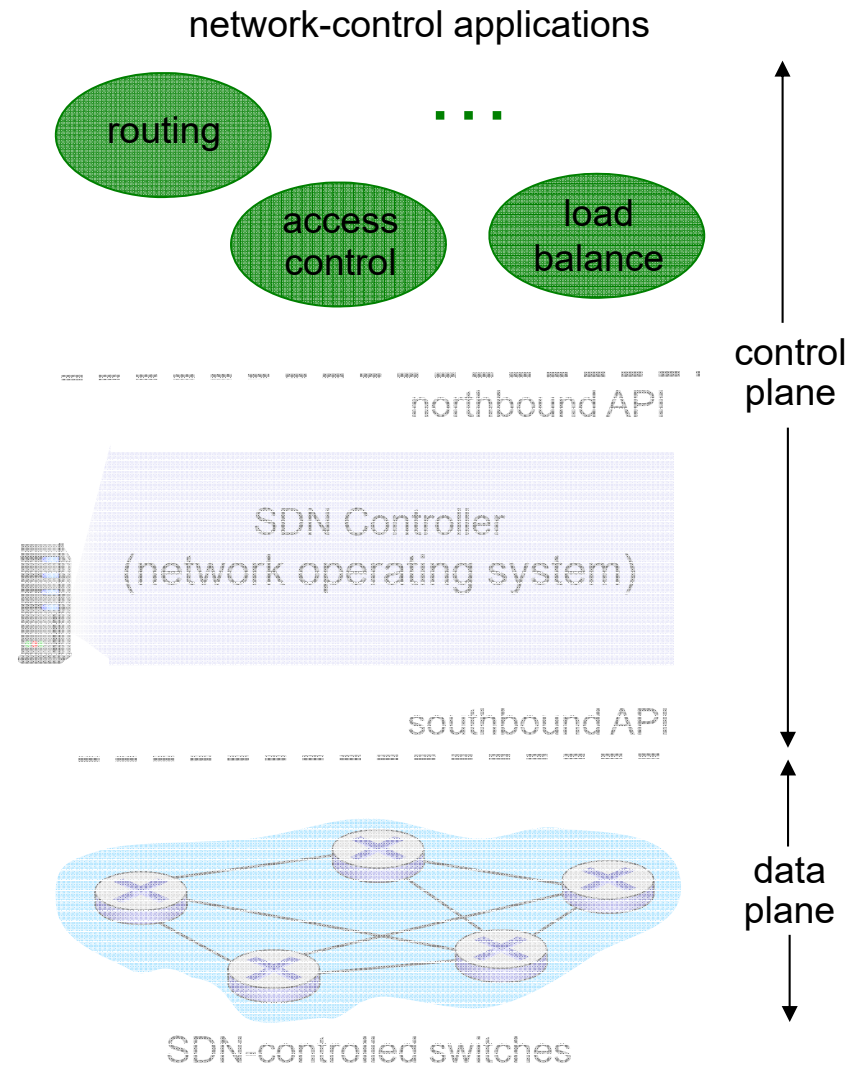
- ❑ Maintain network state information
- ❑ Interacts with network control applications "above" via northbound API
- ❑ Interacts with network switches "below" via southbound API
- ❑ Implemented as distributed system for performance, scalability, fault-tolerance, robustness





# SDN Perspective: Control Applications

- ❑ “Brains” of control:  
implement control functions  
using lower-level services,  
API provided by SDN controller
- ❑ Can be provided by 3rd party
- ❑ Distinct from routing  
vendor or SDN controller



# Chapter 5: Outline

- ❑ 5.1 Introduction
- ❑ 5.2 Routing protocols
  - ❖ Link state
  - ❖ Distance vector
- ❑ 5.3 Intra-As routing in the Internet: OSPF
- ❑ 5.4 Routing among the ISPs: BGP
- ❑ 5.5 The SDN control plane
- ❑ 5.6 ICMP: The Internet Control Message Protocol
- ❑ 5.7 Network management and SNMP

# ICMP: Internet Control Message Protocol

- ❑ Short messages used to send error & other control information
- ❑ Used by hosts & **routers** to communicate network-level information
  - ❖ Error reporting: unreachable host, network, port, protocol
  - ❖ Echo request/reply
    - Used by ping
- ❑ Network-layer "above" IP:
  - ❖ ICMP msgs carried in IP datagrams

Type	Code	description
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	4	frag. needed and DF set
3	6	dest network unknown
3	7	dest host unknown
8	0	echo request (ping)
9	0	route advertisement (mobile)
10	0	router discovery
11	0	TTL expired
12	0	bad IP header

8 bits	8 bits	16 bits
Type	Code	Checksum
Depends on message type		
Internet header + 8 bytes of original datagram causing error		

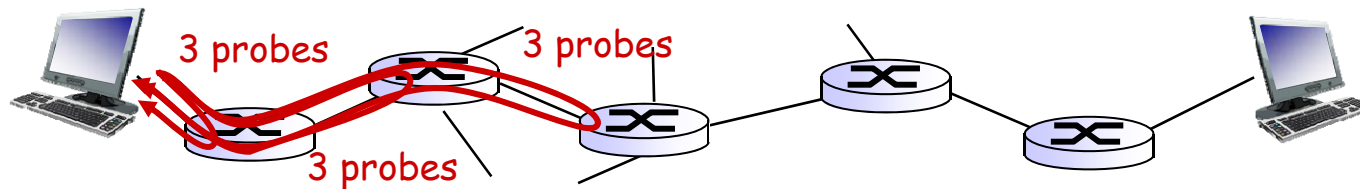
# Traceroute and ICMP

- ❑ Source sends series (3) of
  - ❖ ICMP echo request (**Win**)
  - ❖ UDP segments (**Linux**) to dest with an unlikely port number
    - ❖ 1<sup>st</sup> TTL = 1; 2<sup>nd</sup> TTL = 2, etc.
- ❑ When n<sup>th</sup> datagram arrives at n<sup>th</sup> router:
  - ❖ Router discards datagram
  - ❖ Sends to source an ICMP "TTL expired" message
  - ❖ Message includes name of router & IP address

- ❑ When ICMP message arrives, source calculates RTT
- ❑ Traceroute does this 3 times

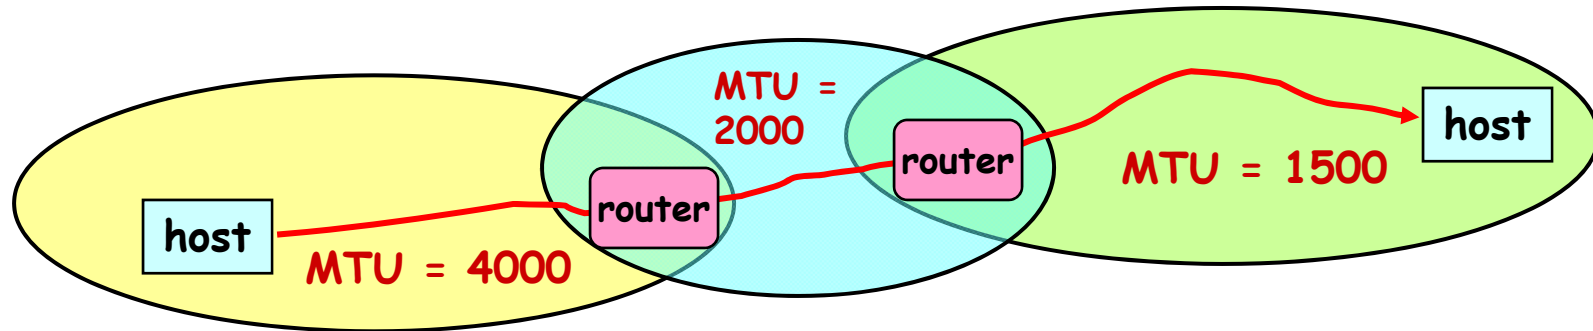
## Stopping criterion

- ❑ **Win**: ICMP echo reply
- ❑ **Linux**: UDP segment eventually arrives at destination host
  - ❖ Destination returns ICMP "dest port unreachable" packet (type 3, code 3)
- ❑ When source gets the ICMP, it stops



# Path MTU Discovery with ICMP

8 bits		16 bits
Type: 3	Code: 4	Checksum
Unused = 0		Next-hop MTU
Internet header + 8 bytes of original datagram		



- ❑ Operation
  - ❖ Send max-sized (MTU) packet with "do not fragment" flag set in IP hdr
  - ❖ If problem encountered, router returns ICMP message
    - "Destination unreachable: Fragmentation needed"
- ❑ Typically, all packets follow same route
  - ❖ Makes sense to do MTU discovery if message is large
    - Send series of packets (i.e., a large message) from one host to another after MTU has been discovered to amortize discovery cost
- ❑ Enabled by default in Windows
  - ❖ To display current MTU
    - `netsh interface ipv4 show subinterfaces`